

Particle Filter vs. Extended Kalman Filter Localization for a Mobile Manipulator

Jacob T. Cassidy*

Johns Hopkins Whiting School of Engineering, Baltimore, Maryland 21218

Reliable localization is a critical prerequisite for mobile manipulation, yet algorithms must frequently contend with accumulating odometry drift and perceptual ambiguity in structured environments. This study presents a comparative evaluation of three estimators: dead reckoning, an Extended Kalman Filter (EKF), and a Monte Carlo Localization Particle Filter (PF). These are evaluated on a Clearpath Husky mobile manipulator equipped with a UR5e arm and Robotiq 2F-85 gripper operating in a simulated ROS 2 factory environment. Relying solely on imperfect wheel odometry and 2D LIDAR scans, the filters are evaluated over a three-waypoint route across a sensitivity grid of varying noise injection levels and particle counts. The objective is to quantify the trade-off between the computational efficiency of the EKF during nominal navigation and the robustness of the Particle Filter when recovering from adversarial multimodal scenarios such as global localization and simulated kidnapping.

I. Overview

Serving as the paper’s introduction, this section will introduce the mobile manipulation platform, a Clearpath Husky A200 carrying a UR5e arm and a Robotiq 2F-85 gripper, equipped with a front-mounted 270° 2D LIDAR and an arch-mounted 360° 2D LIDAR. It will provide a high-level summary of the localization problem in the simulated ROS 2/Gazebo open-source factory environment, explaining how the robot must rely on imperfect wheel odometry and noisy LIDAR scans to determine its pose. Finally, it will state the paper’s core objective: a comparative study testing the claim that a Particle Filter can recover in multimodal scenarios where an EKF fails, while the EKF matches PF accuracy at a lower compute cost during benign runs.

Planned Literature Integration:

- **Macenski, S., et al. (2022) [1]:** This paper will be used for **background** on ROS 2, the middleware that connects the simulator, sensors, and estimators in this study.

II. Problem Statement

Providing necessary background on the state estimation problem, this section will formalize the robot’s localization challenges by identifying specific sources of uncertainty. First, it will describe the unbounded error accumulation in wheel odometry caused by slip, generated by Gazebo’s wheel-slip plugin on the skid-steer base, and by injected odometry noise. Second, it will detail the perceptual ambiguity of LIDAR scans (noise-free in Gazebo, so range noise is injected) in a highly structured, repetitive indoor factory environment. Third, it will explain how unknown initial poses (requiring global localization) and unmodeled displacements (kidnapping via Gazebo teleportation) create multimodal posteriors that violate standard single-Gaussian assumptions.

Planned Literature Integration:

- **Thrun, S., Fox, D., Burgard, W., & Dellaert, F. (2001) [2]:** This paper will be used for **background**, as it characterizes the global localization and kidnapped robot problems, the multimodal posteriors they produce, and the difficulty plain MCL has in recovering from them.

III. Study Approach

Detailing the methodology, this section will outline the mathematical frameworks for the evaluated estimators: dead reckoning, an adapted EKF (fusing odometry with scan-matched pose observations), and a from-scratch Particle Filter (using a likelihood-field LIDAR measurement model). It will explicitly acknowledge the implementation asymmetry

*Graduate Student, Robotics and Autonomous Systems, AIAA Member.

between the adapted EKF and from-scratch PF as a study limitation. It will establish the evaluation protocol: a three-waypoint factory route driven by the simulator’s ground-truth pose, so estimator error never alters the trajectory, with all three estimators processing the same recorded sensor streams, and a sensitivity grid over particle count $N \in \{100, 500, 2000, 5000\}$, odometry motion-noise parameters $\alpha_{1..4} = 0.05k$ with $k \in \{0, 1, 2, 4\}$, and LIDAR range noise $\sigma_r \in \{0.01, 0.03, 0.10\}$ m injected into those streams. Finally, it will justify the 20-trial count per configuration: 20 trials pin each reported average to within about half of the typical run-to-run variation (a 95% confidence interval of roughly ± 0.47 standard deviations), and they separate a filter that reliably recovers from one that reliably fails even with a couple of exceptions, since 18 recoveries in 20 trials bound the true recovery rate above 68% while 2 in 20 bound it below 32% (95% Clopper–Pearson intervals).

Planned Literature Integration:

- **Thrun, S., Burgard, W., & Fox, D. (2005) [3]:** This foundational text will be used for **methodology**, providing the canonical mathematical derivations for both the Extended Kalman Filter and Monte Carlo Localization algorithms.
- **Dellaert, F., Fox, D., Burgard, W., & Thrun, S. (1999) [4]:** This paper introduces the core MCL formulation and will be used for **methodology**, as the formulation the from-scratch Particle Filter follows in localizing against the known map.
- **Leonard, J. J., & Durrant-Whyte, H. F. (1991) [5]:** This paper established EKF localization against an a priori map and will be used for **methodology**, as the basis of the EKF baseline estimator.
- **Gordon, N. J., Salmond, D. J., & Smith, A. F. M. (1993) [6]:** This paper introduced the bootstrap filter and will be used for **methodology**, defining the propagate–weight–resample cycle the Particle Filter implements.
- **Arulampalam, M. S., et al. (2002) [7]:** This tutorial provides a comprehensive overview of sequential Monte Carlo methods and will be used for **background**, specifically justifying the selection of the bootstrap/SIR form for the Particle Filter implementation.

IV. Results

This section will present the quantitative findings from the simulated trials using tables and plots to display position RMSE, heading RMSE, waypoint position error, and per-step compute time across the sensitivity grid. It will feature comparative plots showing mean values with 95% confidence intervals to clearly evaluate the primary hypothesis regarding the accuracy-compute trade-off. Additionally, it will visualize the algorithms’ recovery times in the adversarial tests (global localization from a uniform prior and mid-run kidnapping), measuring the simulation timesteps required for each filter to return below the 0.25 m position error threshold.

Planned Literature Integration:

- **Gutmann, J.-S., & Fox, D. (2002) [8]:** This experimental comparison of Kalman filter, Markov, and Monte Carlo localization in accuracy, robustness, and recovery from robot displacement will be used for **comparison**, as the prior result against which this study’s findings are checked.

A. Preliminary Results

All numbers in this subsection are preliminary. The 20-seed table covers only the nominal operating point; its EKF rows and the sensitivity sweep over N , k and σ_r are still running. Results marked “offline synthetic” come from a 2D test map with simulated odometry and LIDAR, not Gazebo.

Core comparison (Gazebo, 20 seeds). Table 1 reports 20 seeded Gazebo trials per row at the nominal operating point: $N = 2000$ particles, injected odometry noise $k = 1$ ($\alpha_{1..4} = 0.05$) and LIDAR range noise $\sigma_r = 0.03$ m, with the RRT* planner driving the three-waypoint route on ground truth. The EKF rows are being re-run after a fix to how the EKF timestamps its updates, which addresses divergences seen on 2 of the 20 nominal seeds during a fast, roughly 180° in-place turn at the second goal; until then the table shows the PF and dead reckoning only, and the single-run test cases below give the EKF’s behaviour. On the nominal route the PF averaged 0.25 m position error against dead reckoning’s 4.7 m, at a 4.6 ms median update. Under kidnapping it returned within 0.25 m in 16 of 20 trials (95% interval 0.56–0.94) after a mean of 7.8 s. Starting from a uniform prior over the map, it globally localized in 11 of 20 trials (0.32–0.77), taking 55 s on average, so 2000 particles are not enough for reliable global localization in this map; the N sweep will show how the success rate grows with particle count.

Table 1 Core results, Gazebo, 20 seeds per row ($N = 2000$, $k = 1$, $\sigma_r = 0.03$ m). Mean $\pm$ 95% CI; recovery is to within 0.25 m, with its 95% Clopper–Pearson interval. Preliminary; EKF rows pending a re-run.


Scenario	Estimator	Pos. RMSE (m)	Head. RMSE ($^\circ$)	Waypoint err. (m)	Update (ms)	Recovered	Recovery (s)
Nominal	EKF	pending re-run with the EKF timestamp fix					
	PF	0.253 ± 0.072	5.66 ± 2.97	0.133 ± 0.041	4.63 ± 0.07	–	–
	Dead reckoning	4.666 ± 0.495	32.28 ± 4.92	5.987 ± 0.597	–	–	–
Kidnap	EKF	pending re-run with the EKF timestamp fix					
	PF	1.660 ± 0.281	16.27 ± 3.54	2.856 ± 0.473	4.85 ± 0.10	16/20 [0.56, 0.94]	7.8 ± 1.9
Global	EKF	pending re-run with the EKF timestamp fix					
	PF	2.019 ± 1.289	33.70 ± 19.52	2.426 ± 1.859	5.07 ± 0.11	11/20 [0.32, 0.77]	54.9 ± 3.7
	Dead reckoning	4.666 ± 0.495	–	–	–	0/20 [0.00, 0.17]	–

Test cases (Gazebo, single runs). The two test cases for the operable-code deliverable both pass (Figures 1 and 2), with expected outputs derived independently of the filters: the waypoints in the route’s pose file for the nominal case, and the teleport pose for the kidnap case (Table 2). On the nominal route the EKF finished within 6 cm of every waypoint (0.070 m RMSE) and the PF within 11 cm (0.121 m RMSE), against 3.06 m mean error for dead reckoning. After a teleport at $t = 37$ s the PF was back within 0.25 m in 6.1 s, and the EKF never recovered, finishing 21.5 m from the true pose.

Table 2 Test cases, Gazebo, single runs ($N = 2000$, $k = 1$, $\sigma_r = 0.03$ m). Both pass.


Test case	Estimator	Result	Expected output
1: nominal route	EKF	waypoint errors 0.058 / 0.028 / 0.035 m (RMSE 0.070 m)	route waypoints
	PF	waypoint errors 0.075 / 0.111 / 0.059 m (RMSE 0.121 m)	
	Dead reckoning	3.06 m mean error	
2: kidnap at 37 s	EKF	not recovered, 21.5 m final error	teleport pose
	PF	within 0.25 m after 6.1 s	

EKF vs Gazebo ground truth (nominal, $k=1.0$, $\sigma_r=0.03$ m, seed 0)




(a) EKF: 0.035 m final error, 2σ bound 0.037 m.

PF vs Gazebo ground truth (nominal, $k=1.0$, $\sigma_r=0.03$ m, seed 0, $N=2000$)



(b) PF ($N = 2000$): 0.059 m final error, 2σ bound 0.058 m.

Fig. 1 Test case 1 (nominal route, Gazebo, seed 0), final frame of each filter's live view. Left of each: occupancy map, ground truth (green), estimate and 2σ ellipse, and the LIDAR scan projected from the estimate. Right: position and heading error against the filter's own 2σ bound. Both filters stay mostly below the 0.25 m threshold, with brief spikes above it, and their errors repeatedly exceed 2σ , the overconfidence discussed below.



(a) $t = 38.1$ s (1 s after): PF 3.23 m, EKF 8.39 m. (b) $t = 42.1$ s (5 s after): PF 0.13 m, EKF 4.41 m. (c) $t = 52.1$ s (15 s after): PF 0.06 m, EKF 16.63 m.

Fig. 2 Test case 2 (kidnap at $t = 37$ s, Gazebo), intermediate posteriors of both filters at the same moments: PF particles and mean with 2σ ellipse (orange), EKF mean and 2σ ellipse (blue), ground truth (green). The injected particles spread over the map find the true pose within seconds, while the EKF’s single Gaussian stays confidently wrong and drifts further away. At 5 s the PF mean is already close although its 2σ ellipse is still several metres wide, because particles remain spread over the map. Live-view recordings of both test cases are written to `results/test_case_1` and `results/test_case_2` by the test-case launch file.

An earlier Gazebo check replayed two recorded 60 s, roughly 21 m drives through both filters without injected noise (PF at $N = 500$). Odometry alone averaged 3.7–4.0 m of position error; the EKF averaged 0.08 m (0.26 m worst, 0.3° heading) at about 4 ms per scan; the PF averaged 0.15–0.18 m (0.38–0.66 m worst, 0.6 – 0.9° heading) at a 1.25 ms median. Both filters were overconfident: the EKF’s position RMSE was about four times its reported σ , and the PF’s about thirteen times. The final paper will report a covariance consistency check (normalized estimation error squared) alongside RMSE so that this is quantified rather than anecdotal.

Offline synthetic tests. Before the Gazebo runs, the evaluation pipeline was exercised on a synthetic 2D map. On the nominal route the EKF reached 0.03 m and the PF 0.09 m position RMSE. After a kidnap the PF recovered in 11 s and the EKF did not (Figures 3 and 4, which add the error-versus-time traces the Gazebo posterior snapshots lack). Global localization from a uniform prior with 2000 particles converged confidently onto the 180° mirror of the true pose, a symmetry failure that the full study will probe with larger particle counts; it is also a concrete example of the perceptual ambiguity described in the Problem Statement.

EKF vs ground truth (synthetic kidnap, offline test)




Fig. 3 EKF on the offline synthetic kidnap test. Left: map, ground truth (green), EKF estimate and 2σ ellipse (blue), and the LIDAR scan projected from the estimate. Right: position and heading error against the filter's own 2σ bound. After the teleport at about 37 s the error jumps above 10 m and never returns; the covariance grows but the single Gaussian cannot jump to the true pose.

PF vs ground truth (synthetic kidnap, offline test)




Fig. 4 PF on the same offline synthetic kidnap test ($N = 2000$). The error spikes above 15 m at the teleport and falls back under the 0.25 m threshold within about 11 s as injected particles find the true pose.

Still to come. The 20-seed EKF rows and the one-factor-at-a-time sweep over N , k and σ_r will be added here.

V. Conclusion

The final section will synthesize the experimental results to definitively answer whether the primary hypothesis holds true: that the Particle Filter recovers in both adversarial scenarios where the EKF does not, and that the EKF matches or exceeds nominal-route accuracy at lower per-step compute time. It will discuss the practical implications of these findings for the robotic system’s downstream waypoint controller and planning pipeline. Lastly, it will propose future work, such as mitigating the noted implementation asymmetry or extending the evaluation to dynamic obstacles within the factory.

Acknowledgments

Anthropic’s Claude was used to review this outline against the course rubric and the instructor’s feedback on the proposal, to draft revisions to it, and to summarize the preliminary trial results and extract figure stills from the filter visualizations. The localization nodes, evaluation tooling and simulation runs reported here were also developed with Claude’s assistance.

References

- [1] Macenski, S., Foote, T., Gerkey, B., Lalancette, C., and Woodall, W., “Robot Operating System 2: Design, Architecture, and Uses in the Wild,” *Science Robotics*, Vol. 7, No. 66, 2022, p. eabm6074.
- [2] Thrun, S., Fox, D., Burgard, W., and Dellaert, F., “Robust Monte Carlo Localization for Mobile Robots,” *Artificial Intelligence*, Vol. 128, No. 1–2, 2001, pp. 99–141.
- [3] Thrun, S., Burgard, W., and Fox, D., *Probabilistic Robotics*, MIT Press, Cambridge, MA, 2005.
- [4] Dellaert, F., Fox, D., Burgard, W., and Thrun, S., “Monte Carlo Localization for Mobile Robots,” *Proceedings of the IEEE International Conference on Robotics and Automation (ICRA)*, Vol. 2, 1999, pp. 1322–1328.
- [5] Leonard, J. J., and Durrant-Whyte, H. F., “Mobile Robot Localization by Tracking Geometric Beacons,” *IEEE Transactions on Robotics and Automation*, Vol. 7, No. 3, 1991, pp. 376–382.
- [6] Gordon, N. J., Salmond, D. J., and Smith, A. F. M., “Novel Approach to Nonlinear/Non-Gaussian Bayesian State Estimation,” *IEE Proceedings F (Radar and Signal Processing)*, Vol. 140, No. 2, 1993, pp. 107–113.
- [7] Arulampalam, M. S., Maskell, S., Gordon, N., and Clapp, T., “A Tutorial on Particle Filters for Online Nonlinear/Non-Gaussian Bayesian Tracking,” *IEEE Transactions on Signal Processing*, Vol. 50, No. 2, 2002, pp. 174–188.
- [8] Gutmann, J.-S., and Fox, D., “An Experimental Comparison of Localization Methods Continued,” *Proceedings of the IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2002, pp. 454–459.